

PostgreSQL with Django: Complete In-Depth Tutorial

I'll teach you PostgreSQL integration with Django, covering ORM, Migrations, and Docker Integration from beginner to advanced level.

1. Django ORM (Object-Relational Mapper)

What is ORM and Why Do We Need It?

Simple Analogy: Think of ORM as a **translator** between two people who speak different languages. You speak Python, and PostgreSQL speaks SQL. The ORM translates your Python code into SQL queries that PostgreSQL understands.

Without ORM (Raw SQL):

```
python

import psycopg2
cursor.execute("SELECT * FROM tasks WHERE status = 'completed'")
results = cursor.fetchall()
```

With Django ORM:

```
python

tasks = Task.objects.filter(status='completed')
```

Why ORM is Essential:

1. **Security:** Prevents SQL injection attacks
 2. **Portability:** Same code works with different databases
 3. **Maintainability:** Python code is easier to read than complex SQL
 4. **Relationships:** Handles complex table relationships automatically
-

How Django ORM Works Internally

The Three-Layer Architecture

Python Code (models.py)



Django ORM (QuerySet API)



SQL Database (PostgreSQL)

1. Model Definition Phase

python

models.py

from django.db import models

class Task(models.Model):

title = models.CharField(max_length=200)

description = models.TextField()

status = models.CharField(max_length=20, default='pending')

created_at = models.DateTimeField(auto_now_add=True)

updated_at = models.DateTimeField(auto_now=True)

def __str__(self):

return self.title

What Django Does Behind the Scenes:

- Creates a **metaclass** that analyzes your model
- Maps each field to a PostgreSQL column type
- Generates SQL CREATE TABLE statements
- Sets up primary keys, indexes, and constraints

2. QuerySet Lazy Evaluation

Key Concept: Django QuerySets are **lazy** - they don't hit the database until you actually need the data.

python

This doesn't execute any SQL yet!

pending_tasks = Task.objects.filter(status='pending')

SQL is executed only when you iterate or evaluate

for task in pending_tasks: *# NOW SQL executes*

print(task.title)

Django ORM Core Operations

Basic CRUD Operations

1. CREATE Operations

```
python

# Method 1: Create and save
task = Task(title="Learn Django", description="Study ORM concepts")
task.save()

# Method 2: Create in one step
task = Task.objects.create(
    title="Learn PostgreSQL",
    description="Understand database concepts"
)

# Method 3: Bulk create (more efficient for multiple records)
tasks = [
    Task(title="Task 1", description="Description 1"),
    Task(title="Task 2", description="Description 2"),
    Task(title="Task 3", description="Description 3"),
]
Task.objects.bulk_create(tasks)
```

Generated SQL:

```
sql

INSERT INTO myapp_task (title, description, status, created_at, updated_at)
VALUES ('Learn Django', 'Study ORM concepts', 'pending', NOW(), NOW());
```

2. READ Operations

python

Get all records

```
all_tasks = Task.objects.all()
```

Filter records

```
pending_tasks = Task.objects.filter(status='pending')
```

```
completed_tasks = Task.objects.filter(status='completed')
```

Get single record

```
task = Task.objects.get(id=1) # Raises exception if not found
```

```
task = Task.objects.filter(id=1).first() # Returns None if not found
```

Complex filtering

```
recent_tasks = Task.objects.filter(  
    created_at__gte=datetime.now() - timedelta(days=7),  
    status__in=['pending', 'in_progress']  
)
```

Generated SQL Examples:

sql

-- Task.objects.all()

```
SELECT * FROM myapp_task;
```

-- Task.objects.filter(status='pending')

```
SELECT * FROM myapp_task WHERE status = 'pending';
```

-- Complex filter

```
SELECT * FROM myapp_task  
WHERE created_at >= '2024-08-13'  
AND status IN ('pending', 'in_progress');
```

3. UPDATE Operations

```
python
```

```
# Update single record
```

```
task = Task.objects.get(id=1)
```

```
task.status = 'completed'
```

```
task.save()
```

```
# Bulk update (more efficient)
```

```
Task.objects.filter(status='pending').update(status='in_progress')
```

```
# Update with calculations
```

```
from django.db.models import F
```

```
Task.objects.filter(priority__lt=5).update(priority=F('priority') + 1)
```

4. DELETE Operations

```
python
```

```
# Delete single record
```

```
task = Task.objects.get(id=1)
```

```
task.delete()
```

```
# Bulk delete
```

```
Task.objects.filter(status='completed').delete()
```

```
# Delete all (be careful!)
```

```
Task.objects.all().delete()
```

Advanced Querying and Relationships

Model Relationships

python

models.py

```
class User(models.Model):
    username = models.CharField(max_length=150)
    email = models.EmailField()

class Category(models.Model):
    name = models.CharField(max_length=100)

class Task(models.Model):
    title = models.CharField(max_length=200)
    description = models.TextField()
    user = models.ForeignKey(User, on_delete=models.CASCADE) # Many-to-One
    category = models.ForeignKey(Category, on_delete=models.SET_NULL, null=True)
    tags = models.ManyToManyField('Tag') # Many-to-Many

class Tag(models.Model):
    name = models.CharField(max_length=50)
```

Relationship Queries

Forward Relationships

python

Get user of a task

```
task = Task.objects.get(id=1)
user = task.user # This triggers a database query
```

More efficient - use select_related

```
task = Task.objects.select_related('user').get(id=1)
user = task.user # No additional query needed
```

Reverse Relationships

python

Get all tasks for a user

```
user = User.objects.get(id=1)
```

```
user_tasks = user.task_set.all() # Default reverse relation
```

Or with related_name

```
class Task(models.Model):
```

```
    user = models.ForeignKey(User, on_delete=models.CASCADE, related_name='tasks')
```

Now you can use:

```
user_tasks = user.tasks.all()
```

Complex Queries with Joins

python

Get all tasks with their users and categories

```
tasks = Task.objects.select_related('user', 'category').all()
```

Get all users with their tasks

```
users = User.objects.prefetch_related('tasks').all()
```

Complex filtering across relationships

```
tasks = Task.objects.filter(  
    user__username='john_doe',  
    category__name='Work',  
    tags__name='urgent'  
).distinct()
```

Generated SQL:

sql

-- select_related creates JOINS

```
SELECT task.*, user.*, category.*  
FROM myapp_task task  
JOIN myapp_user user ON task.user_id = user.id  
JOIN myapp_category category ON task.category_id = category.id;
```

-- prefetch_related uses separate queries

```
SELECT * FROM myapp_user;  
SELECT * FROM myapp_task WHERE user_id IN (1, 2, 3, ...);
```

Query Optimization Best Practices

1. N+1 Query Problem

Bad Example (Creates N+1 queries):

```
python

tasks = Task.objects.all()
for task in tasks:
    print(task.user.username) # Each iteration hits database!
```

Good Example (Single query with JOIN):

```
python

tasks = Task.objects.select_related('user').all()
for task in tasks:
    print(task.user.username) # No additional queries
```

2. Use only() and defer() for Large Objects

```
python

# Only fetch specific fields
tasks = Task.objects.only('title', 'status').all()

# Defer heavy fields
tasks = Task.objects.defer('description').all()
```

3. Aggregate and Annotate

```
python

from django.db.models import Count, Avg

# Count tasks per user
users = User.objects.annotate(task_count=Count('tasks'))

# Average task completion time
avg_completion = Task.objects.aggregate(Avg('completion_time'))
```

2. Django Migrations

What Are Migrations?

Simple Analogy: Migrations are like **version control for your database schema**. Just like Git tracks changes to your code, migrations track changes to your database structure.

Why Migrations Are Essential:

1. **Team Collaboration:** Everyone gets the same database structure
 2. **Deployment Safety:** Apply changes consistently across environments
 3. **Rollback Capability:** Undo problematic changes
 4. **Change Tracking:** See what changed and when
-

How Django Migration System Works

The Migration Lifecycle

Model Changes → makemigrations → Migration Files → migrate → Database Schema

1. Migration Files Structure

```
python

# migrations/0001_initial.py
from django.db import migrations, models

class Migration(migrations.Migration):
    initial = True

    dependencies = []

    operations = [
        migrations.CreateModel(
            name='Task',
            fields=[
                ('id', models.AutoField(primary_key=True)),
                ('title', models.CharField(max_length=200)),
                ('description', models.TextField()),
                ('status', models.CharField(default='pending', max_length=20)),
                ('created_at', models.DateTimeField(auto_now_add=True)),
            ],
        ),
    ]
```

2. Django Migration Commands Deep Dive

`python manage.py makemigrations`

What it does:

1. Compares current models.py with last migration
2. Detects changes (new models, fields, deletions)
3. Generates Python migration file
4. Assigns sequential number

Example Workflow:

```
bash

# Add new field to model
class Task(models.Model):
    title = models.CharField(max_length=200)
    priority = models.IntegerField(default=1) # NEW FIELD

# Generate migration
python manage.py makemigrations
```

Generated Migration:

```
python

# migrations/0002_task_priority.py
operations = [
    migrations.AddField(
        model_name='task',
        name='priority',
        field=models.IntegerField(default=1),
    ),
]
```

python manage.py migrate

What it does:

1. Checks django_migrations table for applied migrations
2. Applies pending migrations in order
3. Updates django_migrations table
4. Executes SQL commands

Example SQL Generated:

```
sql
```

```
-- For AddField operation
```

```
ALTER TABLE myapp_task ADD COLUMN priority INTEGER DEFAULT 1 NOT NULL;
```

```
-- Django tracks this in:
```

```
INSERT INTO django_migrations (app, name, applied)
```

```
VALUES ('myapp', '0002_task_priority', NOW());
```

Migration Types and Operations

1. Model Operations

```
python
```

```
# Create Model
```

```
migrations.CreateModel(
```

```
    name='Category',
```

```
    fields=[
```

```
        ('id', models.AutoField(primary_key=True)),
```

```
        ('name', models.CharField(max_length=100)),
```

```
    ],
```

```
)
```

```
# Delete Model
```

```
migrations.DeleteModel(name='OldModel'),
```

```
# Rename Model
```

```
migrations.RenameModel(old_name='Task', new_name='TodoItem'),
```

2. Field Operations

python

Add Field

```
migrations.AddField(  
    model_name='task',  
    name='due_date',  
    field=models.DateField(null=True),  
)
```

Remove Field

```
migrations.RemoveField(model_name='task', name='old_field'),
```

Alter Field

```
migrations.AlterField(  
    model_name='task',  
    name='title',  
    field=models.CharField(max_length=300), # Changed from 200  
)
```

Rename Field

```
migrations.RenameField(  
    model_name='task',  
    old_name='desc',  
    new_name='description',  
)
```

3. Data Migrations

Sometimes you need to migrate data, not just schema:

python

```
# migrations/0003_populate_categories.py
```

```
from django.db import migrations
```

```
def populate_categories(apps, schema_editor):
```

```
    Category = apps.get_model('myapp', 'Category')
```

```
    Category.objects.create(name='Work')
```

```
    Category.objects.create(name='Personal')
```

```
    Category.objects.create(name='Shopping')
```

```
def reverse_populate_categories(apps, schema_editor):
```

```
    Category = apps.get_model('myapp', 'Category')
```

```
    Category.objects.filter(name__in=['Work', 'Personal', 'Shopping']).delete()
```

```
class Migration(migrations.Migration):
```

```
    dependencies = [
```

```
        ('myapp', '0002_category'),
```

```
    ]
```

```
    operations = [
```

```
        migrations.RunPython(populate_categories, reverse_populate_categories),
```

```
    ]
```

Advanced Migration Techniques

1. Migration Dependencies

python

```
class Migration(migrations.Migration):
```

```
    dependencies = [
```

```
        ('myapp', '0001_initial'),
```

```
        ('otherapp', '0003_add_user_profile'), # Cross-app dependency
```

```
    ]
```

2. Custom Migration Operations

```
python
```

```
# Custom SQL
```

```
migrations.RunSQL(  
    "CREATE INDEX idx_task_status_created ON myapp_task(status, created_at);",  
    reverse_sql="DROP INDEX idx_task_status_created;"  
)
```

```
# Custom Python code
```

```
migrations.RunPython(  
    code=forward_func,  
    reverse_code=reverse_func  
)
```

3. Squashing Migrations

When you have many small migrations, combine them:

```
bash
```

```
python manage.py squashmigrations myapp 0001 0005
```

Migration Best Practices and Common Issues

1. Safe Migration Practices

```
python
```

```
# ❌ Dangerous - removes data
```

```
migrations.RemoveField(model_name='task', name='important_data')
```

```
# ✅ Safe - add new field first, migrate data, then remove old
```

```
# Step 1: Add new field
```

```
migrations.AddField(  
    model_name='task',  
    name='new_field',  
    field=models.CharField(max_length=100, null=True)  
)
```

```
# Step 2: Data migration (separate migration)
```

```
# Step 3: Remove old field (another migration)
```

2. Handling Migration Conflicts

```
bash
```

```
# When migrations conflict
```

```
python manage.py makemigrations --merge
```

3. Rolling Back Migrations

```
bash
```

```
# Rollback to specific migration
```

```
python manage.py migrate myapp 0001
```

```
# Rollback all migrations for an app
```

```
python manage.py migrate myapp zero
```

4. Production Migration Strategy

```
python
```

```
# settings.py - Different migration settings for production
```

```
if ENVIRONMENT == 'production':
```

```
    # Require explicit migration approval
```

```
    MIGRATION_ATOMIC = True
```

```
    MIGRATION_TIMEOUT = 300
```

3. Docker Integration with PostgreSQL

What is Docker and Why Use It?

Simple Analogy: Docker is like a **shipping container** for your application. Just like shipping containers standardize how goods are transported, Docker containers standardize how applications run across different environments.

Benefits for Django + PostgreSQL:

1. **Consistency:** Same environment everywhere (dev, test, prod)
2. **Isolation:** No conflicts with system PostgreSQL
3. **Easy Setup:** New developers get running quickly
4. **Version Control:** Lock specific PostgreSQL versions

Setting Up PostgreSQL with Docker Compose

1. Basic docker-compose.yml Structure

yaml

docker-compose.yml

version: '3.8'

services:

PostgreSQL Database Service

db:

image: postgres:15

container_name: django_postgres

restart: always

environment:

POSTGRES_DB: myproject_db

POSTGRES_USER: myproject_user

POSTGRES_PASSWORD: myproject_password

ports:

- "5432:5432"

volumes:

- postgres_data:/var/lib/postgresql/data

- ./init.sql:/docker-entrypoint-initdb.d/init.sql

networks:

- django_network

Django Web Application Service

web:

build: .

container_name: django_web

restart: always

command: python manage.py runserver 0.0.0.0:8000

ports:

- "8000:8000"

volumes:

- ./app

environment:

- DEBUG=1

- DATABASE_URL=postgresql://myproject_user:myproject_password@db:5432/myproject_db

depends_on:

- db

networks:

- django_network

Named volumes for data persistence

volumes:

postgres_data:

Custom network for service communication


```
networks:  
  django_network:  
    driver: bridge
```

2. Understanding Each Component

PostgreSQL Service Configuration

```
yaml  
  
db:  
  image: postgres:15 # Official PostgreSQL image, version 15  
  container_name: django_postgres # Custom container name  
  restart: always # Restart container if it crashes
```

Environment Variables:

- `POSTGRES_DB`: Database name to create
- `POSTGRES_USER`: Database user to create
- `POSTGRES_PASSWORD`: Password for the user

Volume Mounting Explained

```
yaml  
  
volumes:  
  - postgres_data:/var/lib/postgresql/data # Named volume for persistence  
  - ./init.sql:/docker-entrypoint-initdb.d/init.sql # Initialization script
```

Volume Types:

1. **Named Volume** (`postgres_data`): Managed by Docker, persists data
2. **Bind Mount** (`./init.sql`): Links host file to container

Network Configuration

```
yaml  
  
networks:  
  django_network:  
    driver: bridge
```

This creates an isolated network where services can communicate using service names as hostnames.

Django Configuration for Docker PostgreSQL

1. Environment-Based Configuration

```
python

# settings.py
import os
from pathlib import Path

# Build paths inside the project like this: BASE_DIR / 'subdir'.
BASE_DIR = Path(__file__).resolve().parent.parent

# Database Configuration
DATABASES = {
    'default': {
        'ENGINE': 'django.db.backends.postgresql',
        'NAME': os.environ.get('DB_NAME', 'myproject_db'),
        'USER': os.environ.get('DB_USER', 'myproject_user'),
        'PASSWORD': os.environ.get('DB_PASSWORD', 'myproject_password'),
        'HOST': os.environ.get('DB_HOST', 'db'), # 'db' is the service name
        'PORT': os.environ.get('DB_PORT', '5432'),
    }
}

# Alternative: Using DATABASE_URL
import dj_database_url
DATABASES = {
    'default': dj_database_url.parse(
        os.environ.get('DATABASE_URL', 'postgresql://myproject_user:myproject_password@db:5432/myproject_db')
    )
}
```

2. Environment Variable Management

.env File (for development)

```
bash

# .env
DEBUG=1
SECRET_KEY=your-secret-key-here
DB_NAME=myproject_db
DB_USER=myproject_user
DB_PASSWORD=myproject_password
DB_HOST=db
DB_PORT=5432
DATABASE_URL=postgresql://myproject_user:myproject_password@db:5432/myproject_db
```

Loading Environment Variables

```
python
```

```
# settings.py
```

```
from dotenv import load_dotenv
```

```
import os
```

```
# Load environment variables from .env file
```

```
load_dotenv()
```

```
# Now use os.environ.get() as shown above
```

3. Requirements and Dockerfile

requirements.txt

```
txt
```

```
Django>=4.2.0
```

```
psycopg2-binary>=2.9.0
```

```
python-dotenv>=1.0.0
```

```
dj-database-url>=2.0.0
```

Dockerfile

dockerfile

Dockerfile

FROM python:3.11-slim

Set environment variables

ENV PYTHONDONTWRITEBYTECODE 1

ENV PYTHONUNBUFFERED 1

Set work directory

WORKDIR /app

Install system dependencies

RUN apt-get update \
 && apt-get install -y --no-install-recommends \
 postgresql-client \
 build-essential \
 libpq-dev \
 && rm -rf /var/lib/apt/lists/*

Install Python dependencies

COPY requirements.txt /app/

RUN pip install --no-cache-dir -r requirements.txt

Copy project

COPY . /app/

Expose port

EXPOSE 8000

Run migrations and start server

CMD ["python", "manage.py", "runserver", "0.0.0.0:8000"]

Advanced Docker Configuration

1. Multi-Environment Setup

Development Environment

yaml

```
# docker-compose.dev.yml
```

```
version: '3.8'
```

```
services:
```

```
  db:
```

```
    image: postgres:15
```

```
    environment:
```

```
      POSTGRES_DB: myproject_dev
```

```
      POSTGRES_USER: dev_user
```

```
      POSTGRES_PASSWORD: dev_password
```

```
    volumes:
```

```
      - postgres_dev_data:/var/lib/postgresql/data
```

```
    ports:
```

```
      - "5433:5432" # Different port for dev
```

```
  web:
```

```
    build: .
```

```
    command: python manage.py runserver 0.0.0.0:8000
```

```
    volumes:
```

```
      - ./app
```

```
    environment:
```

```
      - DEBUG=1
```

```
      - DATABASE_URL=postgresql://dev_user:dev_password@db:5432/myproject_dev
```

```
    ports:
```

```
      - "8000:8000"
```

```
    depends_on:
```

```
      - db
```

```
volumes:
```

```
  postgres_dev_data:
```

Production Environment

yaml

```
# docker-compose.prod.yml
version: '3.8'
services:
  db:
    image: postgres:15
    environment:
      POSTGRES_DB: myproject_prod
      POSTGRES_USER: prod_user
      POSTGRES_PASSWORD_FILE: /run/secrets/db_password
    secrets:
      - db_password
    volumes:
      - postgres_prod_data:/var/lib/postgresql/data
    networks:
      - internal

  web:
    build:
      context: .
      dockerfile: Dockerfile.prod
    command: gunicorn myproject.wsgi:application --bind 0.0.0.0:8000
    environment:
      - DEBUG=0
      - DATABASE_URL_FILE=/run/secrets/database_url
    secrets:
      - database_url
    depends_on:
      - db
    networks:
      - internal
      - web

secrets:
  db_password:
    file: ./secrets/db_password.txt
  database_url:
    file: ./secrets/database_url.txt

volumes:
  postgres_prod_data:

networks:
  internal:
  web:
```

external: true

2. Health Checks and Waiting

```
yaml

# docker-compose.yml with health checks
services:
  db:
    image: postgres:15
    healthcheck:
      test: ["CMD-SHELL", "pg_isready -U myproject_user -d myproject_db"]
      interval: 30s
      timeout: 10s
      retries: 3
      start_period: 60s
    # ... other configurations

  web:
    # ... other configurations
    depends_on:
      db:
        condition: service_healthy
```

3. Initialization Scripts

```
sql

-- init.sql
-- This runs when the container starts for the first time

-- Create additional databases
CREATE DATABASE myproject_test;

-- Create extensions
CREATE EXTENSION IF NOT EXISTS "uuid-osspl";
CREATE EXTENSION IF NOT EXISTS "pg_trgm";

-- Create custom indexes or functions
-- ...
```

Docker Workflow and Best Practices

1. Development Workflow

```
bash
```

```
# Start services
```

```
docker-compose up -d
```

```
# Run migrations
```

```
docker-compose exec web python manage.py migrate
```

```
# Create superuser
```

```
docker-compose exec web python manage.py createsuperuser
```

```
# View logs
```

```
docker-compose logs -f web
```

```
docker-compose logs -f db
```

```
# Stop services
```

```
docker-compose down
```

```
# Stop and remove volumes (careful - deletes data!)
```

```
docker-compose down -v
```

2. Data Management

Backup Database

```
bash
```

```
# Create backup
```

```
docker-compose exec db pg_dump -U myproject_user myproject_db > backup.sql
```

```
# Restore backup
```

```
docker-compose exec -T db psql -U myproject_user myproject_db < backup.sql
```

Access Database Directly

```
bash
```

```
# Connect to PostgreSQL container
```

```
docker-compose exec db psql -U myproject_user -d myproject_db
```

```
# Or from outside container
```

```
docker-compose exec db psql -U myproject_user -d myproject_db -c "SELECT * FROM myapp_task;"
```

3. Environment Separation Best Practices


```
bash
```

```
# Use different compose files for different environments
```

```
docker-compose -f docker-compose.dev.yml up # Development
```

```
docker-compose -f docker-compose.prod.yml up # Production
```

```
# Override specific services
```

```
docker-compose -f docker-compose.yml -f docker-compose.override.yml up
```

4. Security Best Practices

1. Use secrets for passwords:

```
yaml
```

```
secrets:
```

```
  postgres_password:
```

```
    file: ./secrets/postgres_password.txt
```

2. Limit network exposure:

```
yaml
```

```
services:
```

```
  db:
```

```
    # Don't expose port in production
```

```
    # ports: - "5432:5432" # Remove this line
```

3. Use non-root user:

```
dockerfile
```

```
RUN adduser --disabled-password --gecos '' appuser
```

```
USER appuser
```

4. Set resource limits:

```
yaml
```

```
services:
```

```
  db:
```

```
    deploy:
```

```
      resources:
```

```
        limits:
```

```
          memory: 512M
```

```
          cpus: '0.5'
```

Troubleshooting Common Issues

1. Connection Issues

Problem: `django.db.utils.OperationalError: could not connect to server`

Solutions:

```
bash
```

```
# Check if containers are running
```

```
docker-compose ps
```

```
# Check container logs
```

```
docker-compose logs db
```

```
docker-compose logs web
```

```
# Verify network connectivity
```

```
docker-compose exec web ping db
```

```
# Check PostgreSQL is accepting connections
```

```
docker-compose exec db pg_isready -U myproject_user
```

2. Permission Issues

Problem: `permission denied for database`

Solution:

```
sql
```

```
-- Connect as superuser and grant permissions
```

```
GRANT ALL PRIVILEGES ON DATABASE myproject_db TO myproject_user;
```

```
GRANT ALL PRIVILEGES ON ALL TABLES IN SCHEMA public TO myproject_user;
```

3. Migration Issues

Problem: `relation "django_migrations" does not exist`

Solution:

```
bash

# Ensure database is created and accessible
docker-compose exec web python manage.py migrate --run-syncdb
```

4. Data Persistence Issues

Problem: Data disappears when containers restart

Solution:

```
yaml

# Ensure named volumes are used
volumes:
  - postgres_data:/var/lib/postgresql/data # Not ./data
```

Summary and Next Steps

You've now learned:

1. **Django ORM:** How to interact with PostgreSQL through Python objects, understanding query generation, relationships, and optimization
2. **Migrations:** How Django tracks and applies database schema changes safely across environments
3. **Docker Integration:** How to containerize your Django + PostgreSQL stack for consistent development and deployment

Key Takeaways:

- Always use `select_related()` and `prefetch_related()` for relationship queries
- Test migrations thoroughly before applying to production
- Use named volumes for data persistence in Docker
- Separate configuration for different environments
- Always backup your data before major migrations

Practice Project: Create a task management system with users, categories, and tags using all these concepts together!